Cafy Agent — Standard Operating Procedures (SOPs)

Nov 7, 2026 · Cafy Agent v2.1

How to use these SOPs

Cafy Agent is a UK-focused career platform: it reads your résumé, matches it to live UK job listings, writes a tailored résumé for a chosen job without inventing facts, and follows up recruiter emails with drafts you approve. Each SOP below covers one area; read SOP 1 first, then the ones you need. Every step names the screen, the field to fill and the button to press, using the exact labels in the app.

Reading the screenshots

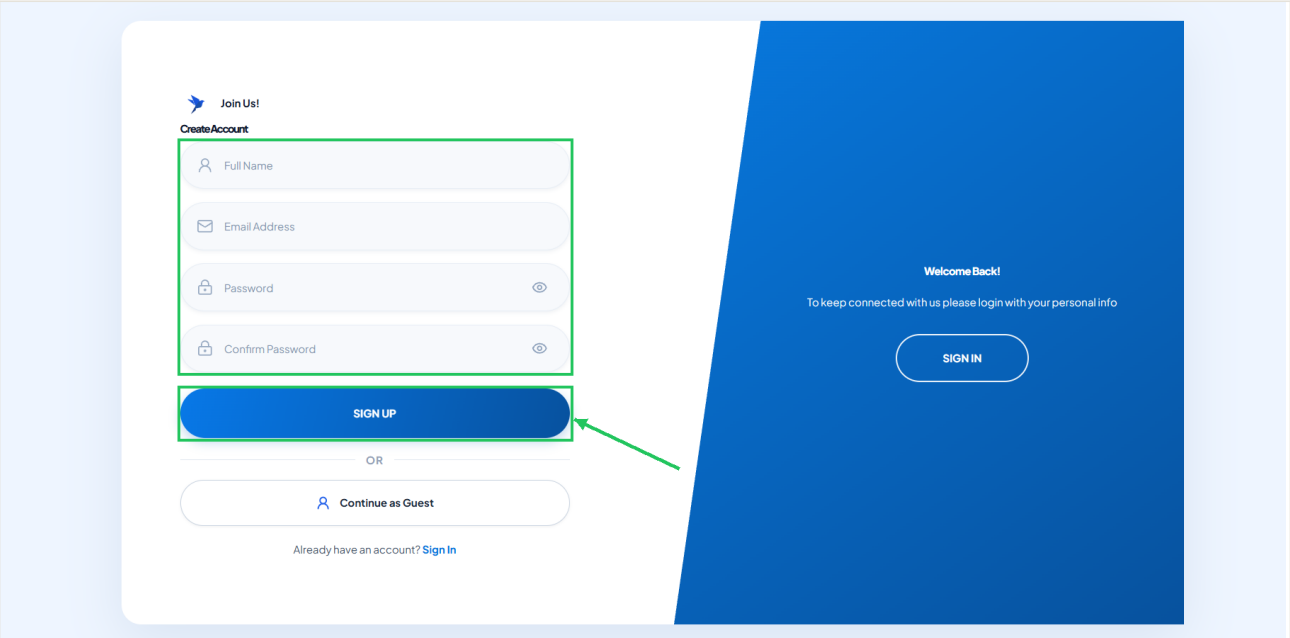
- **Green box** = the field or filter you fill or pick.
- **Green arrow** = the button you press to save or move on.
- Numbers in the steps match the order you work in on the screen.

Contents

Section	What it covers
Signing in, Golden rules, Data flow	Every user — read first
SOP 1 — Candidate: résumé and job search	Upload, matches, filters, job details, visa information
SOP 2 — Candidate: tailored résumé	Create, edit, preview, export, versions
SOP 3 — Candidate: applications	Pipeline, Applied To, Apply now, export
SOP 4 — Email Agent and Google Calendar	Connect Gmail, Google Calendar, Outlook; recruiter emails, approvals, interviews, follow-ups
SOP 5 — Career Analyst and Analytics	AI chat, market insights
SOP 6 — Guest user	What a guest can and cannot do
SOP 7 — Administrator / Operator	Install with uv, configuration, scraping, deployment, backup
How matching, tailoring and sponsorship work	The rules behind the numbers
Troubleshooting, limitations, API summary	Fixes and reference

Signing in (every user)

1. Open Cafy Agent in the browser (computer, tablet or phone).
2. New user: on **Create Account** fill **Full Name**, **Email Address**, **Password** (8–128 characters) and **Confirm Password**, then press **SIGN UP**.
3. Returning user: press **Sign In**, fill **Email Address** and **Password**, press **SIGN IN**.
4. Where the server allows it, **Continue as Guest** starts a private temporary session (see SOP 6).
5. To finish, press **Log Out** at the bottom of the sidebar.



Screen 0.1 — Create Account: fill the boxed fields, press SIGN UP.

After 5 wrong passwords for one email (or 30 from one network) in 15 minutes the app shows *"Too many failed attempts. Try again in N minute(s)."* Wait, then sign in again. Sessions renew automatically while you work.

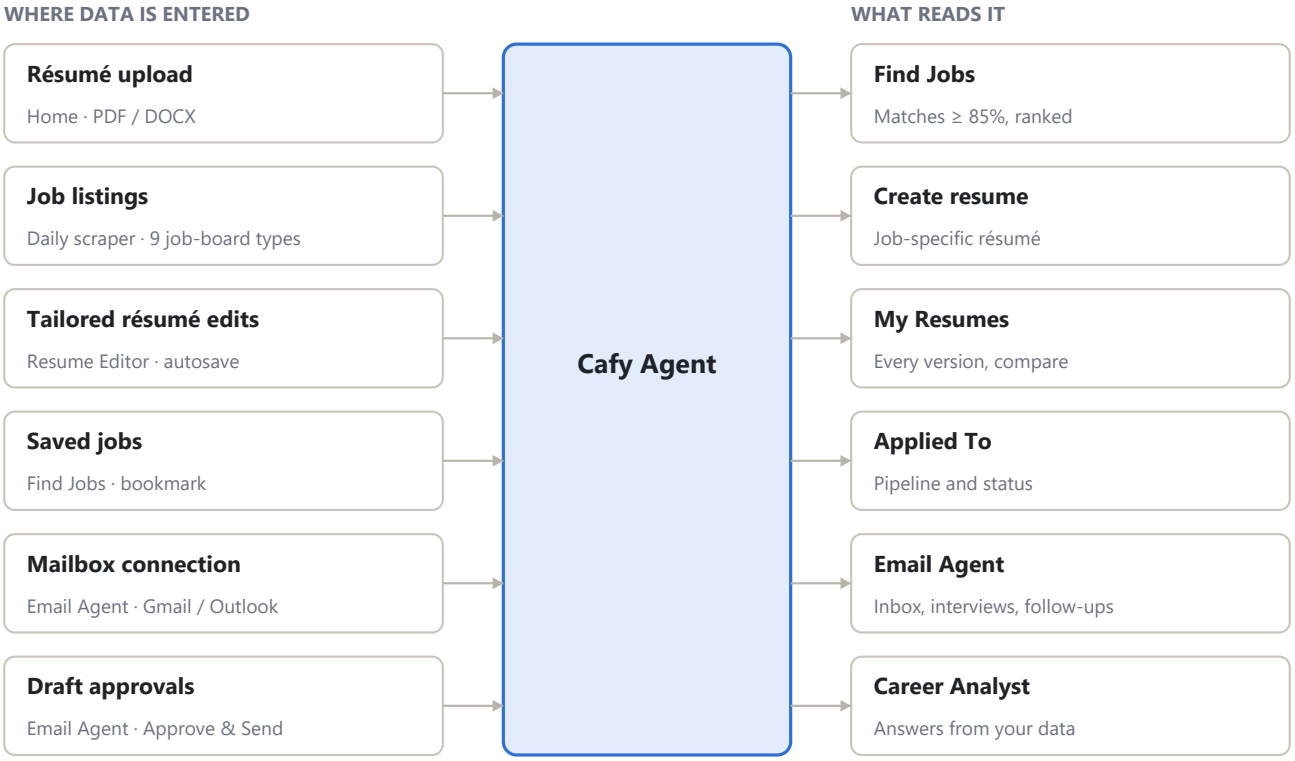
Golden rules

- Upload your **real, current résumé**. Every match score, tailored résumé and email draft reads from it.
- Cafy never invents experience: a skill or keyword is only added to a tailored résumé when your résumé already shows it.
- Nothing is sent to a recruiter until you press **Approve & Send**. Cafy never submits job applications for you.
- Each account sees only its own résumé, versions, emails and applications.
- Edits in the Resume Editor save automatically two seconds after you stop typing.

Data flow — what happens to the data you enter

Every record is entered once and every other screen reads that same record; nobody re-types it. Your résumé is parsed once on upload and drives matching, tailoring, analysis and the career chat. Jobs are collected daily from employer job boards and checked against the UK sponsor register.

Every input is stored once and read by every module



Data flow · 6 entry points, 6 screens that read them

The table below names, for each entry, exactly where it shows up next.

When you enter…	Who	It immediately shows in…	Later it drives…
A résumé (upload)	Candidate	Home Active Profile card; Find Jobs list of matches	Every match score, tailored résumé, Career Analyst answers, data export
Filters and sort	Candidate	Find Jobs list and job count	Nothing else (view only)
A job saved (bookmark / Save all)	Candidate	Applied To; sidebar count	Email Agent application list (on its next load)
Create resume / Regenerate	Candidate	Resume Editor and PDF preview	My Resumes versions; downloads; export ZIP
Editor changes	Candidate	PDF preview; match score refresh	Saved version used for downloads
A mailbox connection	Candidate	Email Agent → Connected accounts	Inbox Intelligence, application statuses, interview and follow-up drafts
Approve & Send	Candidate	Email History → Sent	Follow-up timing
A Google Calendar connection	Candidate	Email Agent → Google Calendar card: Connected	Free interview slots, interview events and follow-up reminders
Daily job scrape	System (scheduler)	Find Jobs; Analytics	New matches for every candidate

SOP 1 — Candidate: résumé and job search

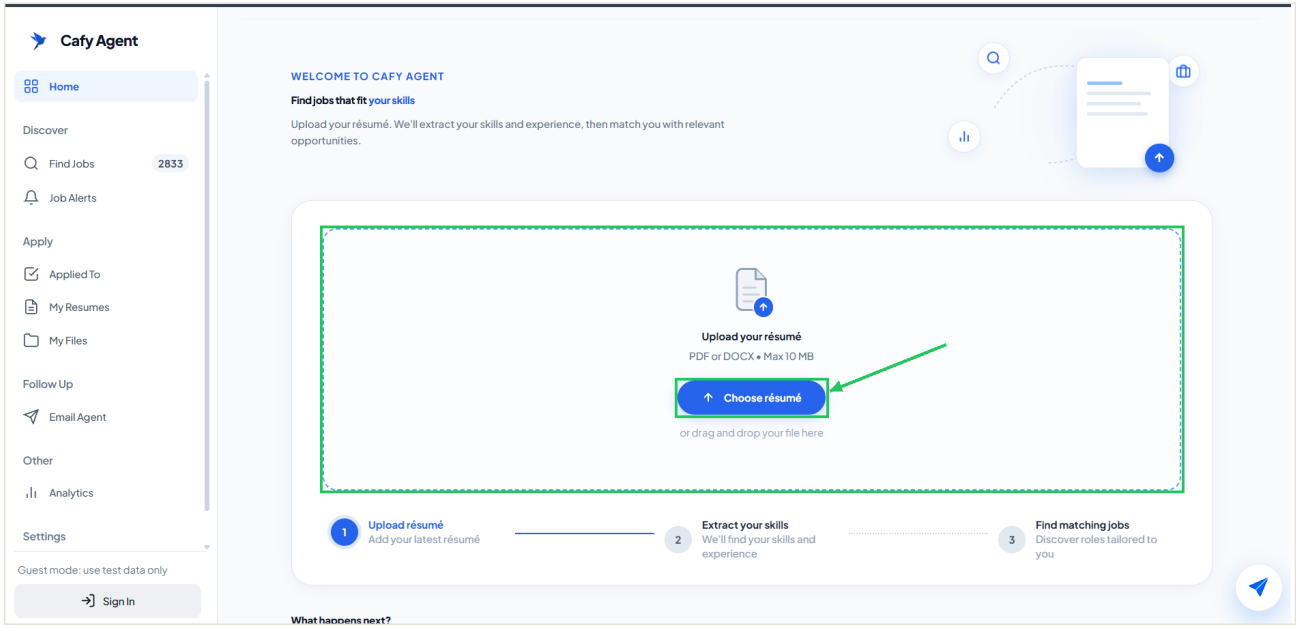
The candidate uploads one résumé; Cafy extracts skills, experience and education and lists only the UK jobs that genuinely fit. You land on **Home**.

Daily routine

When	Screen	What you do
Once	Home	Upload your résumé (replace it when it changes)
Daily	Find Jobs	Review new matches, open details, save good ones
Per job	Create resume	Generate and check a tailored résumé (SOP 2)
Daily	Email Agent	Approve replies, confirm interviews (SOP 4)
Weekly	Applied To, My Resumes	Update statuses; tidy versions

1.1 Upload your résumé (Home)

1. Open **Home** from the sidebar.
2. Press **Choose résumé** (or drag the file onto the box). Accepted: PDF, DOCX or plain text (TXT); maximum 10 MB. Old binary .doc files are not read correctly — save as DOCX or PDF first.
3. Wait while the button shows *Processing Résumé…*. The toast *"Analyzed résumé. Found N matching roles."* appears.
4. Cafy opens **Find Jobs** automatically with your matches.



Screen 1.1 — Home: press Choose résumé (arrow) or drop the file in the box.

After upload, Home shows the **Active Profile** card: file name, name, target role and years, your first five skills, and the buttons **View N Jobs** and **Replace**. Use **Replace** to upload a newer résumé.

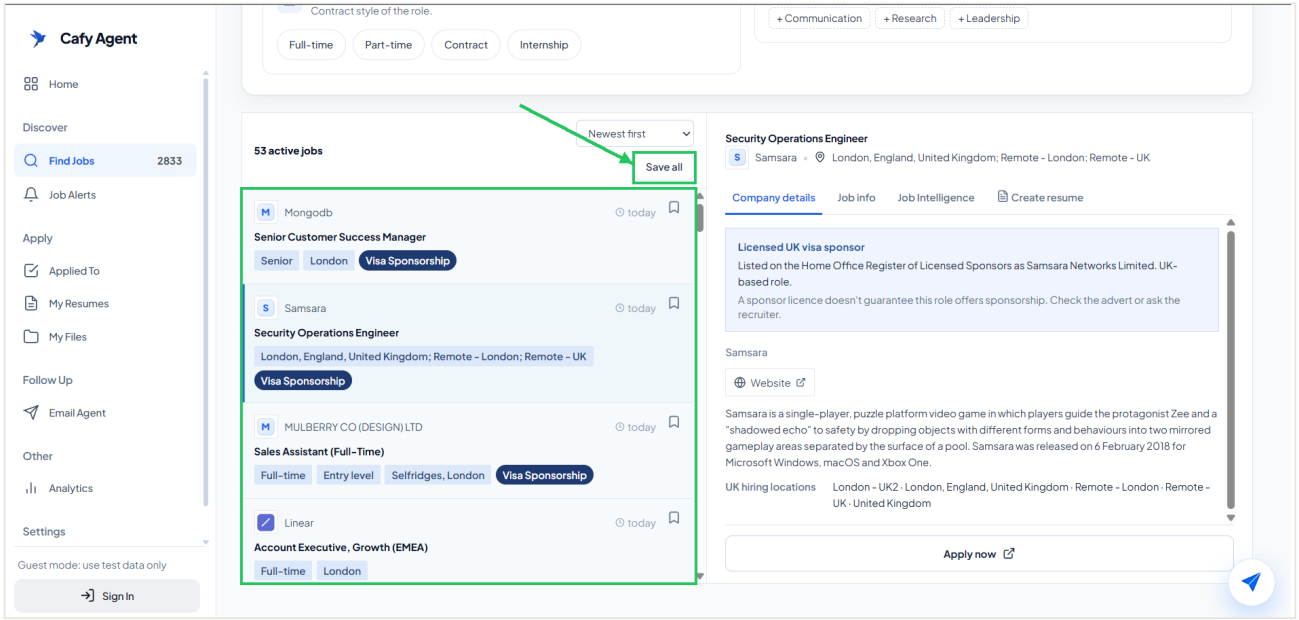
What Cafy reads from the résumé

- Contact details (email, phone, LinkedIn, GitHub, portfolio, location) and name.

- Sections by heading: summary, skills (e.g. *Skills, Core Skills, Technical Skills, Skills & Tools*), experience, internships, education, projects, certifications, publications, achievements.
- Years of experience from the dates of full-time roles, or the years stated in the résumé when roles have no dates (internships never count as years), and a career level: **fresher** (under 1 year) or **experienced**.

1.2 Review your matches (Find Jobs)

1. Open **Find Jobs**. With a résumé loaded, the note *"You're seeing jobs matched to your résumé and preferences."* appears and only qualifying jobs are listed (see *How matching works*).
2. Each row shows company, posting age, title, **N% Match**, tags and **Visa Sponsorship** when the employer is a licensed sponsor and the job is confirmed in the UK.
3. Change the order with **Sort by**: *Newest first*, *Best match* or *Company A–Z*. Press the refresh icon to reload matches.
4. Search with **Job title** and **Company, city or "UK"** at the top; the list filters as you type.

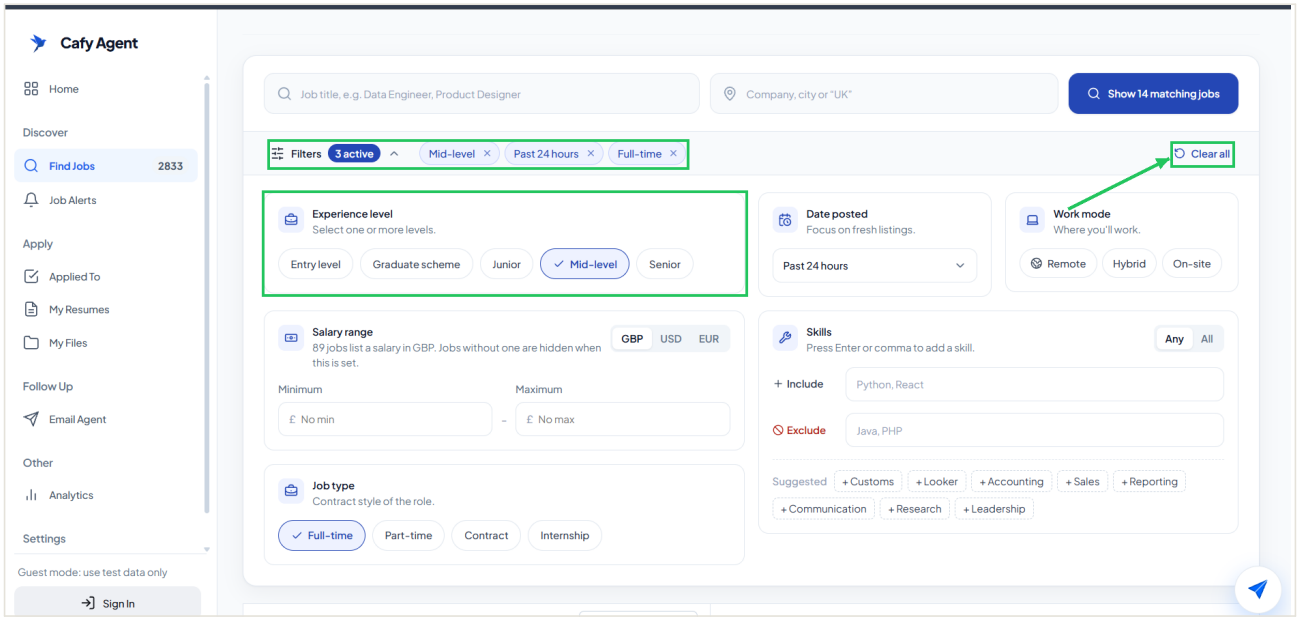


Screen 1.2 — Find Jobs: the job list (box); Save all adds every visible job to your pipeline (arrow). With a résumé loaded, rows also show N% Match and the toolbar shows Remove résumé.

If the list shows *"No roles match your filter criteria"*, no current job meets the matching rules or your filters are too narrow. Clear filters first; if it stays empty, your résumé genuinely has no strong match among today's jobs.

1.3 Filter the list

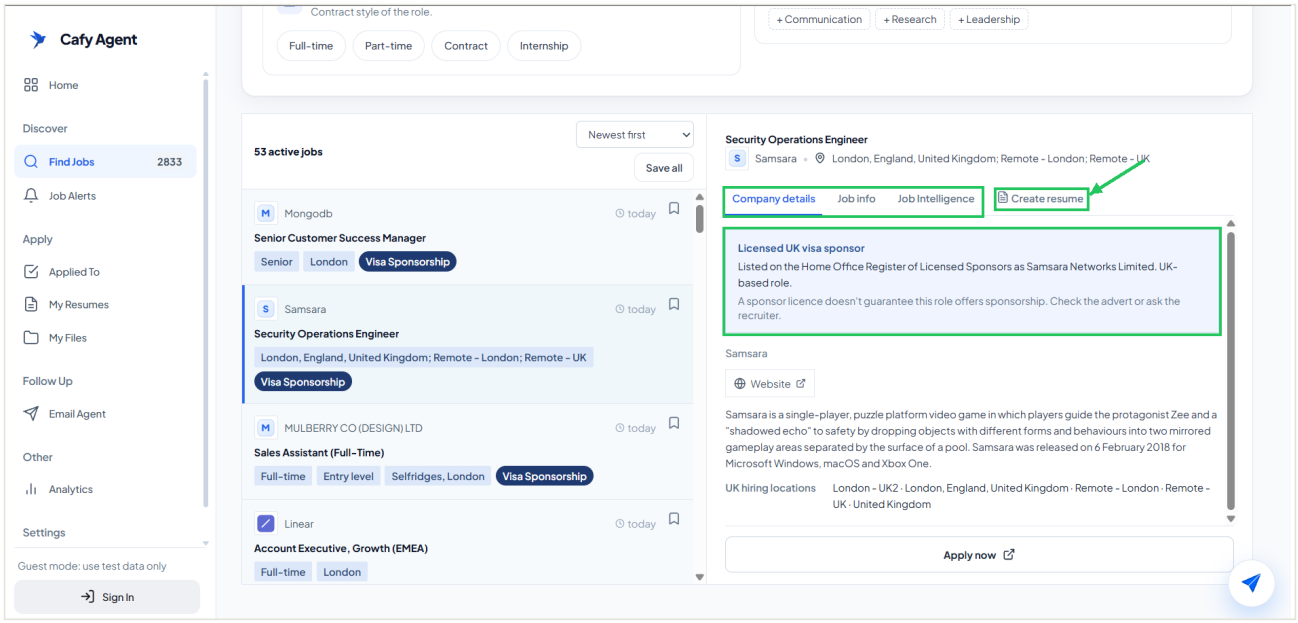
1. Press **Filters**. The panel opens only when you click it; active filters stay visible as chips.
2. Pick any of: **Experience level** (Entry level, Graduate scheme, Junior, Mid-level, Senior), **Date posted** (Any time … Past month), **Work mode** (Remote, Hybrid, On-site), **Salary range** (GBP/USD/EUR, Minimum, Maximum), **Skills** (Include / Exclude, Any or All), **Job type** (Full-time, Part-time, Contract, Internship).
3. Remove one filter with the × on its chip, or all with **Clear all**.



Screen 1.3 — Filters open: the Filters bar with active chips and a filter group (boxes); Clear all removes every filter (arrow).

1.4 Read a job (detail pane)

1. Click a job in the list. The detail pane opens with tabs **Company details**, **Job info** and **Job Intelligence**.
2. **Company details**: the visa callout (e.g. *Licensed UK visa sponsor*, *Not on the sponsor register*), website and LinkedIn, industry, headquarters, size, careers page and UK hiring locations.
3. **Job info**: type, level, mode, salary, posting date and the description. If the description is short, use **Read the full job advert**.
4. **Job Intelligence**: *About the role*, *Hiring process* (each stage marked Confirmed by employer, Reported by candidates or Estimated), *How you match* and *Sources*. Use **Refresh** for new research and **Report an error** if something is wrong.
5. Press **Apply now** to open the employer's advert, or **Create resume** to tailor your résumé (SOP 2).



Screen 1.4 — Job detail: tabs and the visa callout (boxes); Create resume (arrow).

1.5 Save jobs to your pipeline

1. Press the bookmark icon on a job row (toast *"Added job to pipeline!"*), or **Save all** to add every visible job.
2. The sidebar **Applied To** count goes up. Saved jobs are kept for the current browser session only; reloading the page clears them.

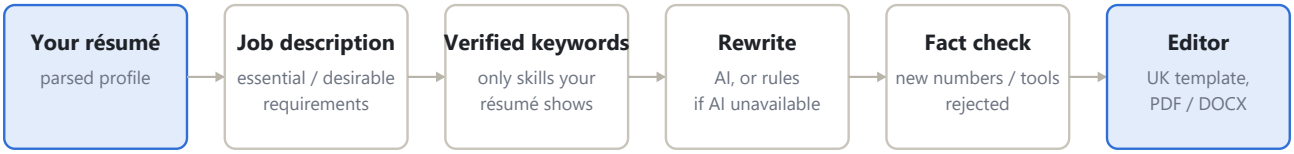
1.6 Remove your résumé

1. In **Find Jobs** press **Remove résumé** in the list toolbar.
2. Confirm *"Remove your saved résumé and uploaded file?"*. The profile, file and match list are cleared and Find Jobs shows all jobs again.

SOP 2 — Candidate: tailored résumé

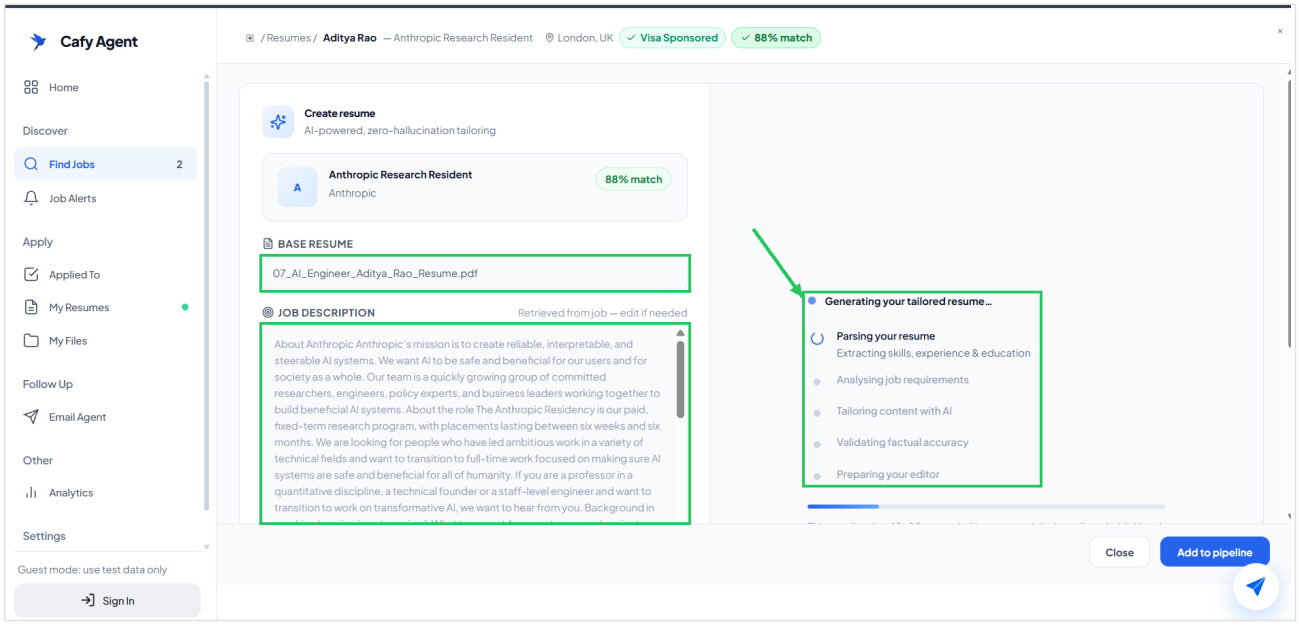
A tailored résumé is built for one selected job from the facts already in your résumé. Cafy rewrites wording, orders content by relevance and adds job keywords only where your résumé proves them.

Tailoring pipeline — from your résumé to a job-specific version



2.1 Create a tailored résumé

1. In **Find Jobs** open a job and press **Create resume**. Generation starts straight away; no review step is needed.
2. On the **Create resume** screen you can check **Base resume**, edit **Job description** (*Retrieved from job — edit if needed*) and add **Additional instructions** (optional, up to 500 characters, e.g. *Emphasise my Python experience*).
3. Press **Generate Tailored Resume** (or **Regenerate Tailored Resume**). Progress shows for 10–20 seconds.
4. The **Resume Editor** opens with score tiles *Job Match*, *Skills Match*, *ATS Readiness* and *Content Quality*.



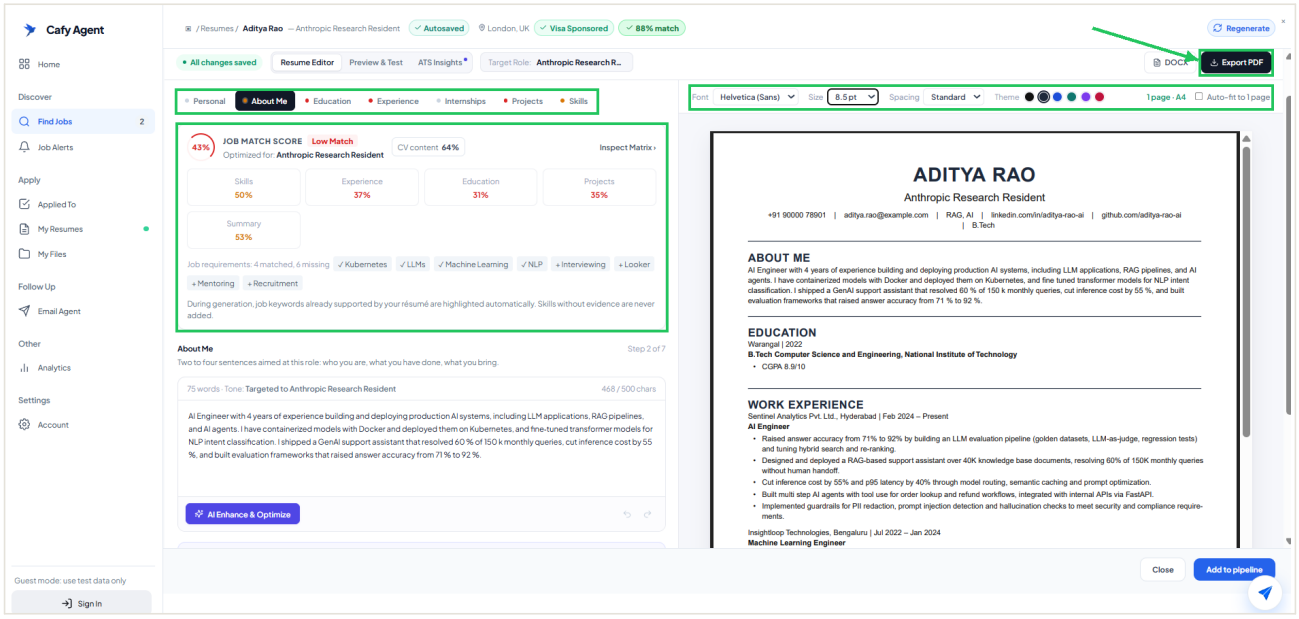
Screen 2.1 — Create resume: Base resume and Job description (boxes); generation progress on the right (arrow).

What the generator does — and never does

- Writes **About Me** for the selected job from real facts: degree or role, internships, skills shown in your work, one real result, and the target role and company.
- Adds a job keyword to Skills only when your résumé shows that skill; essential job skills come first. Always check added skills (see Known limitations).
- Puts your most job-relevant bullets and projects first. Employers, titles, dates and education are copied unchanged.
- Never adds a skill, number, employer or qualification you do not have. Essential skills you lack are listed in the result as *not claimed*.
- If the AI service is unavailable, the résumé is still produced with rule-based tailoring, labelled as such.

2.2 Edit the résumé (Resume Editor)

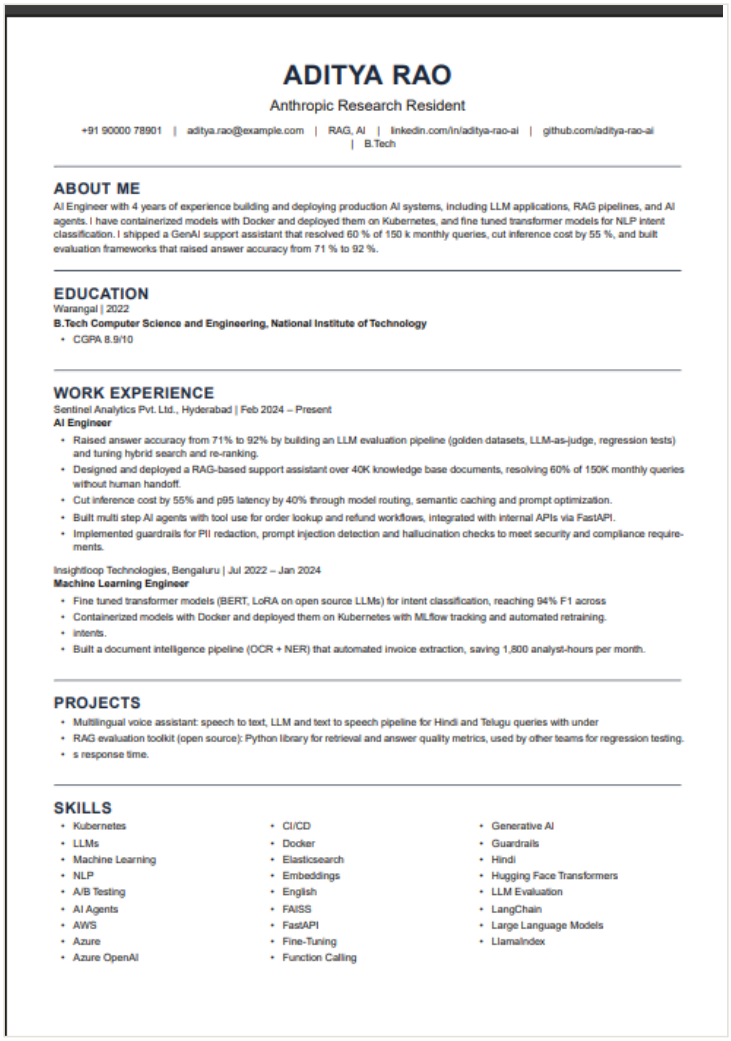
1. Use the section tabs: **Personal**, **About Me**, **Education**, **Experience**, **Internships**, **Projects**, **Skills**. Each shows *Step X of 7*.
2. Type directly in the fields. **Improve with AI** (or the sparkle on a bullet) polishes the wording of text you wrote; it is instructed never to add experience you have not entered — review the result.
3. Add entries with **+ Add … Entry** and bullets with **+ Add Bullet Point**; drag the handle to reorder.
4. In **Skills**, type a skill and press Enter (or paste a comma list). Under *Skills this job asks for*, add only those that are true for you.
5. The status shows *Saving…* then *All changes saved*. Ctrl+S saves now; Ctrl+Z / Ctrl+Y undo and redo.



Screen 2.2 — Resume Editor: section tabs, JOB MATCH SCORE card and PDF settings (boxes); Export PDF (arrow).

2.3 Preview, layout and download

1. The preview on the right uses the UK template: name, role, contact line, then **About Me**, **Education**, **Work Experience** (full-time and internships), **Projects** and **Skills** in three columns.
2. Adjust **Font**, **Size** (8.5–11.5 pt), **Spacing** (Compact, Standard, Relaxed) and **Theme**. Tick **Auto-fit to 1 page** to shrink (never below 9.5 pt). The indicator shows *1 page · A4*.
3. Press **Export PDF** to download the PDF, or **DOCX** for an editable Word file.



ADITYA RAO

Anthropic Research Resident

+91 90000 78901 | aditya.rao@example.com | RAG, AI | linkedin.com/in/aditya-rao-ai | github.com/aditya-rao-ai | B.Tech

ABOUT ME

AI Engineer with 4 years of experience building and deploying production AI systems, including LLM applications, RAG pipelines, and AI agents. I have containerized models with Docker and deployed them on Kubernetes, and fine tuned transformer models for NLP intent classification. I shipped a GenAI support assistant that resolved 60 % of 150 k monthly queries, cut inference cost by 55 %, and built evaluation frameworks that raised answer accuracy from 71 % to 92 %.

EDUCATION

Warangal | 2022

B.Tech Computer Science and Engineering, National Institute of Technology

- CGPA 8.9/10

WORK EXPERIENCE

Sentinel Analytics Pvt. Ltd., Hyderabad | Feb 2024 – Present

AI Engineer

- Raised answer accuracy from 71% to 92% by building an LLM evaluation pipeline (golden datasets, LLM-as-judge, regression tests) and tuning hybrid search and re-ranking.
- Designed and deployed a RAG-based support assistant over 40K knowledge base documents, resolving 60% of 150K monthly queries without human handoff.
- Cut inference cost by 55% and p95 latency by 40% through model routing, semantic caching and prompt optimization.
- Built multi step AI agents with tool use for order lookup and refund workflows, integrated with internal APIs via FastAPI.
- Implemented guardrails for PII redaction, prompt injection detection and hallucination checks to meet security and compliance requirements.

Insightloop Technologies, Bengaluru | Jul 2022 – Jan 2024

Machine Learning Engineer

- Fine tuned transformer models (BERT, LoRA on open source LLMs) for intent classification, reaching 94% F1 across
- Containerized models with Docker and deployed them on Kubernetes with MLflow tracking and automated retraining.
- intents.
- Built a document intelligence pipeline (OCR + NER) that automated invoice extraction, saving 1,800 analyst-hours per month.

PROJECTS

- Multilingual voice assistant: speech to text, LLM and text to speech pipeline for Hindi and Telugu queries with under
- RAG evaluation toolkit (open source): Python library for retrieval and answer quality metrics, used by other teams for regression testing.
- s response time.

SKILLS

- Kubernetes
- LLMs
- Machine Learning
- NLP
- A/B Testing
- AI Agents
- AWS
- Azure
- Azure OpenAI
- CI/CD
- Docker
- Elasticsearch
- Embeddings
- English
- FAISS
- FastAPI
- Fine-Tuning
- Function Calling
- Generative AI
- Guardrails
- Hindi
- Hugging Face Transformers
- LLM Evaluation
- LangChain
- Large Language Models
- LlamaIndex

Screen 2.3 — Exported PDF in the UK template: About Me, Education, Work Experience, Projects, Skills in three columns.

Certifications, publications and achievements stay stored but are not printed in this template.

2.4 Check the match (ATS Insights)

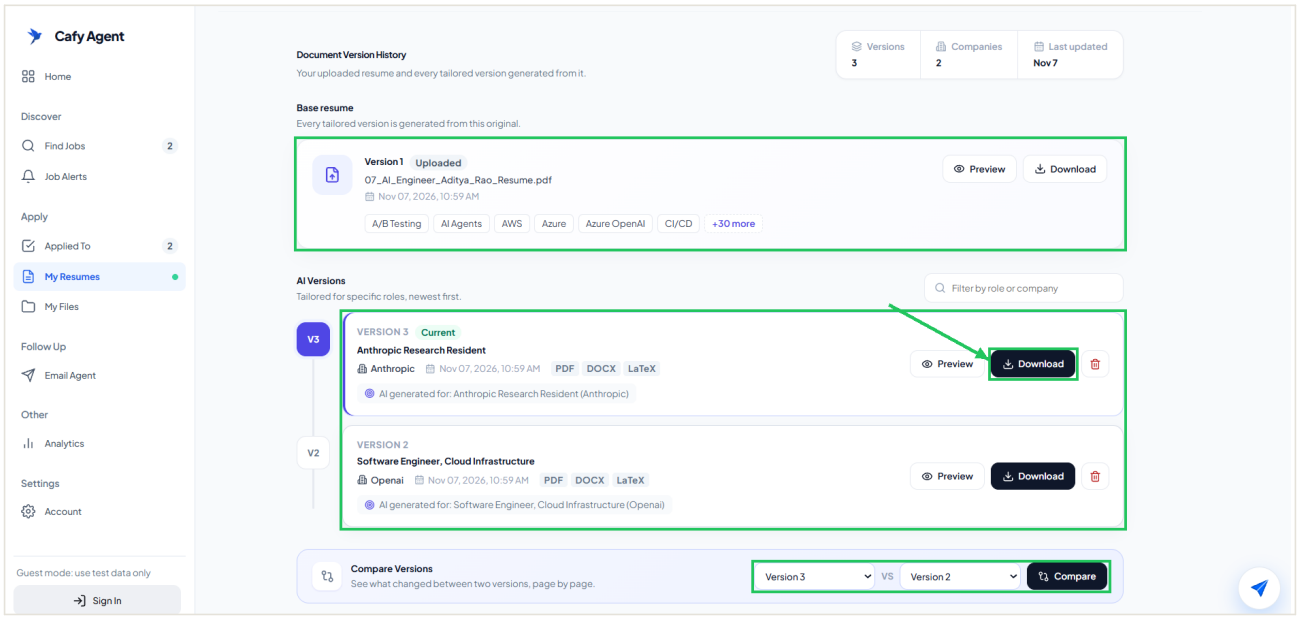
1. The **JOB MATCH SCORE** card above the editor updates as you edit (*Scoring your resume against this job...*). Chips show Skills, Experience, Education, Projects and Summary.
2. Press **Inspect Matrix ›** or open the **ATS Insights** tab for *Actionable Recommendations* such as *Strengthen your opening summary* and *Role keyword coverage*.
3. Act on a recommendation only where it is true for you, then re-check the score.

Screenshot to be added
Screen 2.4

Screen 2.4 — ATS Insights: recommendations (box), Improve with AI (arrow).

2.5 Manage versions (My Resumes)

1. Open **My Resumes**. **Base resume** holds your uploaded original: **Preview** or **Download**.
2. **AI Versions** lists every tailored version, newest first, with the role and company. Use **Filter by role or company**.
3. **Preview** opens a version; **Download** gives PDF (else DOCX, else LaTeX).
4. To compare, pick two versions under **Compare Versions** and press **Compare** (tabs *Page view* and *Changes*).
5. To delete, press the bin icon and confirm. This cannot be undone.

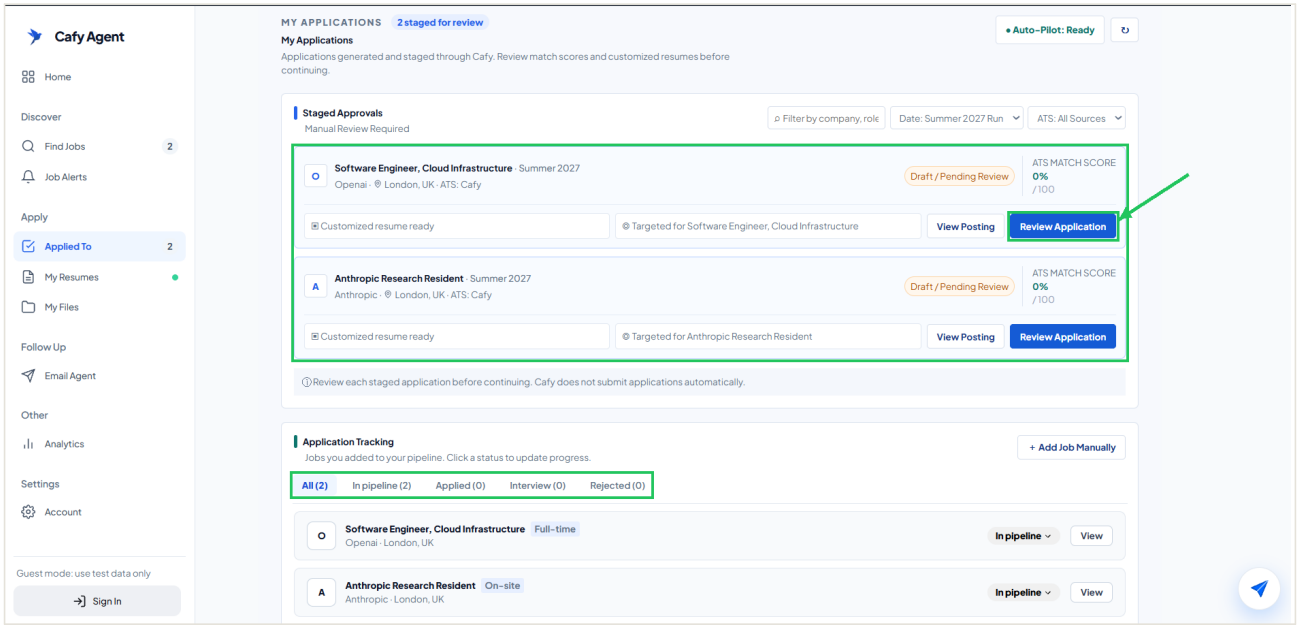


Screen 2.5 — My Resumes: base résumé, AI versions and Compare Versions (boxes); Download (arrow).

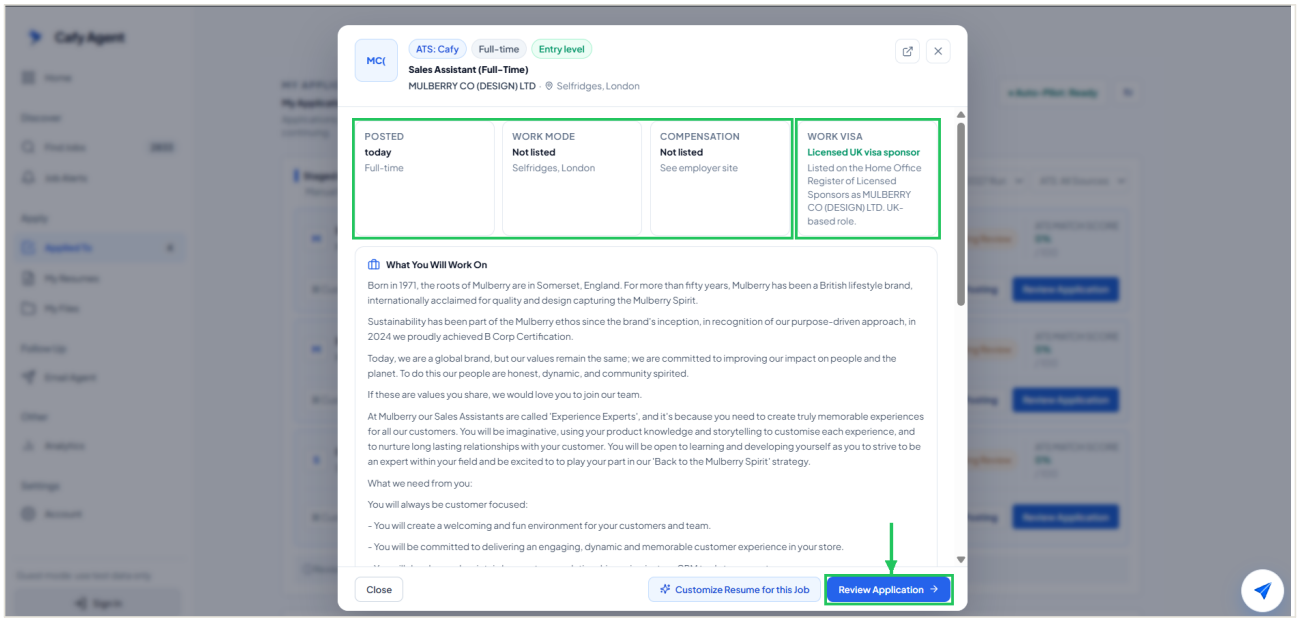
SOP 3 — Candidate: applications and your data

3.1 Track applications (Applied To)

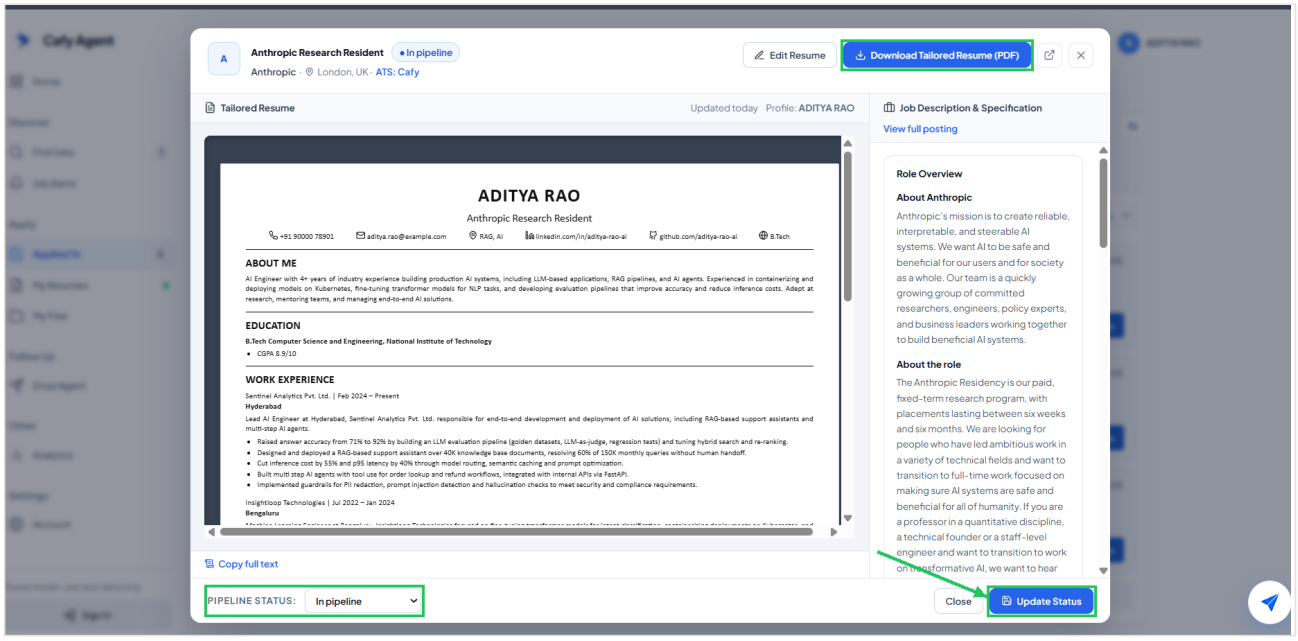
1. Open **Applied To** (page heading *My Applications*).
2. **Staged Approvals** shows the first saved jobs with **View Posting** and **Review Application**. Cafy does not submit applications automatically.
3. In **Application Tracking**, filter with the tabs *All*, *In pipeline*, *Applied*, *Interview*, *Rejected*. Click a status pill to move it: In pipeline → Applied → Interview → Rejected.
4. **Review Application** opens the tailored résumé next to the job: **Edit Resume** / **Create Resume**, **Download Tailored Resume**, **Copy full text**, then set **PIPELINE STATUS** and press **Update Status**.



Screen 3.1 — Applied To: Staged Approvals and status tabs (boxes); Review Application (arrow).



Screen 3.2 — View Posting: posting facts and Work visa status (boxes); Review Application (arrow). Customize Resume for this Job opens Create resume.



Screen 3.3 — Review Application: Download Tailored Resume and PIPELINE STATUS (boxes); Update Status (arrow).

Statuses and saved jobs live in the browser session; reloading the page clears them.

3.2 Apply to a job

1. From the editor or job pane press **Apply now**. Cafy records the application and opens the employer's page in a new tab.
2. If you already applied through Cafy, it asks *"You already applied to this job through Cafy (date). Apply again anyway?"*.
3. Complete and submit the form on the employer's site yourself.

3.3 Export your data (My Files)

1. Open **My Files**. It shows the current file and the number of skills extracted.
2. Under **Export Your Data** tick **Profile Data**, **Resumes** (PDF / DOCX / LaTeX), **Application History**, or **Complete Data** for all three.
3. Press **Create Export**. A ZIP downloads (*"Export ready: filename"*).

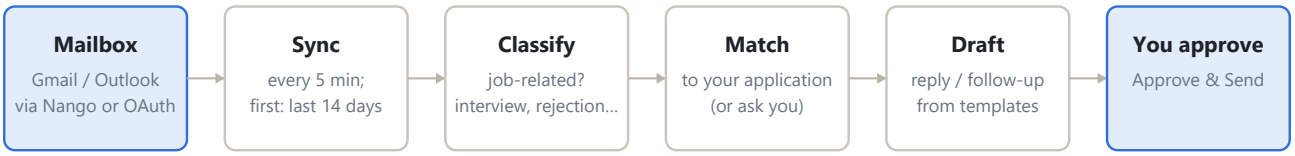
Screenshot to be added
Screen 3.4

Screen 3.4 — My Files: pick what to export (boxes), Create Export (arrow).

SOP 4 — Email Agent and Google Calendar

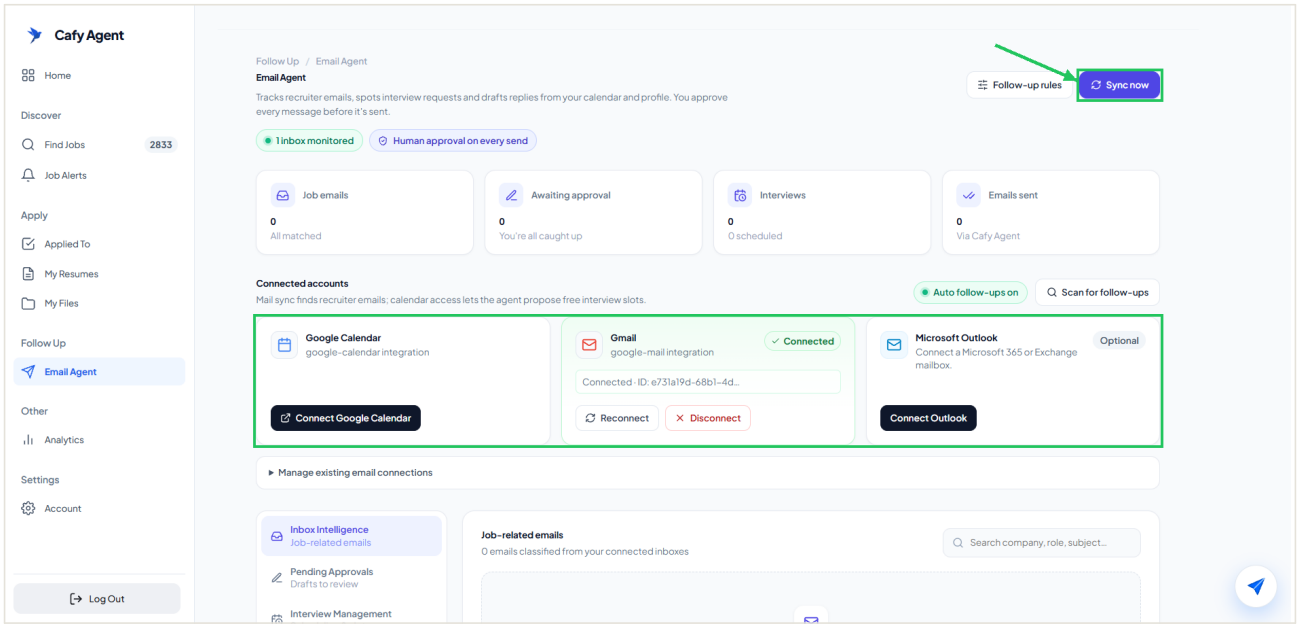
The Email Agent reads recruiter emails, links them to your applications, updates statuses, spots interview invitations and drafts replies and follow-ups. You approve every message before it is sent.

Email pipeline — from recruiter email to an approved reply



4.1 Connect a mailbox

1. Open **Email Agent**. Under **Connected accounts** press **Connect Gmail** (and **Connect Google Calendar** for interview slots). Complete the Google popup.
2. For Microsoft 365 / Exchange press **Connect Outlook** and sign in.
3. The first sync reads the last 14 days in the background; afterwards new mail is checked every 5 minutes (on a Vercel deployment: once a day by cron). Press **Sync now** any time.
4. If a card says *Nango is not configured* or Outlook is disabled, ask the administrator (SOP 7).



Screen 4.1 — Email Agent: Connect Gmail / Calendar / Outlook (boxes), Sync now (arrow).

4.2 Review job emails (Inbox Intelligence)

1. Press **Sync now** to check immediately (*"Synced: N job-related email(s)…"*).
2. In **Inbox Intelligence**, read job emails by category (application received, under review, assessment, interview, rejection, offer…). Personal and promotional mail is ignored; its content is not stored.
3. Under *Confirm which application these belong to*, pick the application or **Not job-related** when Cafy is unsure.

4. Application statuses update automatically only when the classification is confident; a status never moves backwards.

4.3 Approve or reject drafts (Pending Approvals)

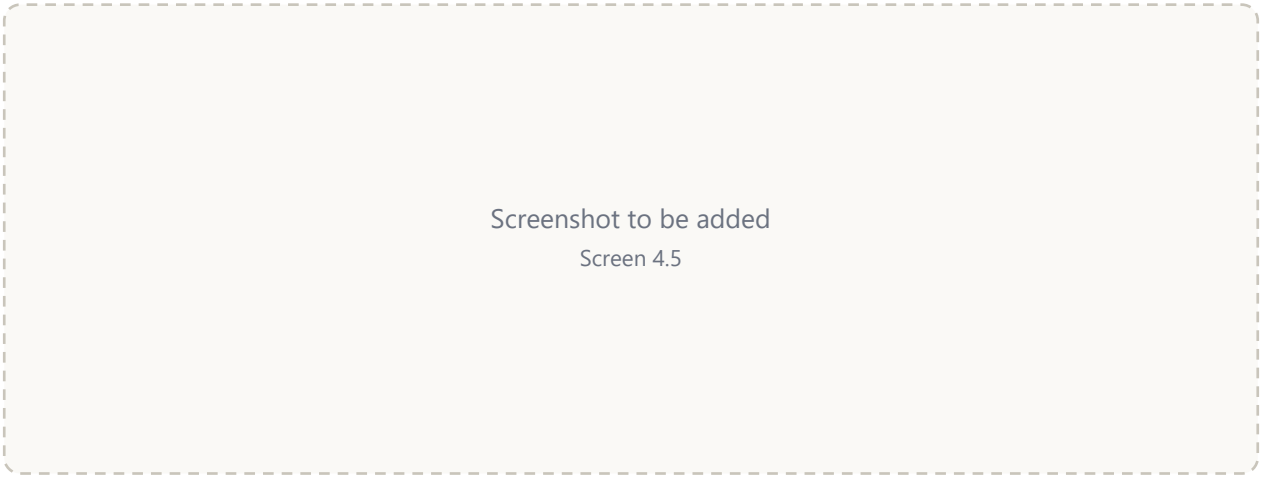
1. Open **Pending Approvals**. Each draft shows **To**, **Subject** and the message body.
2. Edit freely, then press **Save edits**, **Approve & Send**, or **Reject**. Nothing is sent until you approve.

4.4 Schedule interviews (Interview Management)

1. Detected interviews show date, time, duration, format and meeting link.
2. Pick slots under *Free slots from your calendar* (weekdays 09:00–12:00 and 13:00–17:00), then press **Draft reply with N slots**, or **Draft acceptance**.
3. When agreed, fill the date, time, duration and link and press **Confirm & add to calendar**. Cafy re-checks your calendar and adds reminders 1 day and 1 hour before.

4.5 Follow-ups

1. Press **Scan for follow-ups** (it also runs automatically). Rules: wait at least 7 days (or the reply time the employer promised), 7 days between follow-ups, at most 2, and never after the recruiter has replied.
2. Each follow-up is a draft in **Pending Approvals** plus a 15-minute calendar reminder; use **Add to calendar** in **Email History**.



Screen 4.5 — Pending Approvals: edit the draft (box), Approve & Send (arrow).

4.6 Google Calendar

Google Calendar lets the Email Agent offer interview times you are actually free for, put confirmed interviews in your calendar and remind you about follow-ups. It is connected on the Email Agent page, next to Gmail.

Calendar is used to…	What Cafy does	Where you see it
Find free interview slots	Reads busy times only (free/busy) for the next 7 working days; offers up to 6 slots, weekdays 09:00–12:00 and 13:00–17:00	Interview Management → Free slots from your calendar

Calendar is used to…	What Cafy does	Where you see it
Check for clashes	Re-checks the time before confirming; flags an interview that conflicts with an event	Interview Management; Needs your attention
Create interview events	Adds the event with reminders 1 day and 1 hour before	Your Google Calendar
Remind you to follow up	Adds a 15-minute reminder at 10:00 on the next weekday	Email History → Add to calendar

Connect Google Calendar

1. Open **Email Agent**. Under **Connected accounts**, find the **Google Calendar** card (*google-calendar integration*).
2. Press **Connect Google Calendar** and complete the Google sign-in popup, allowing calendar access.
3. The card turns green and shows **Connected · ID**. Connect **Gmail** as well: calendar actions work together with a connected mailbox.
4. To stop access, press **Disconnect** on the card and confirm. Cafy then uses the calendar access of the connected Gmail account, if it has any.

Use it for an interview

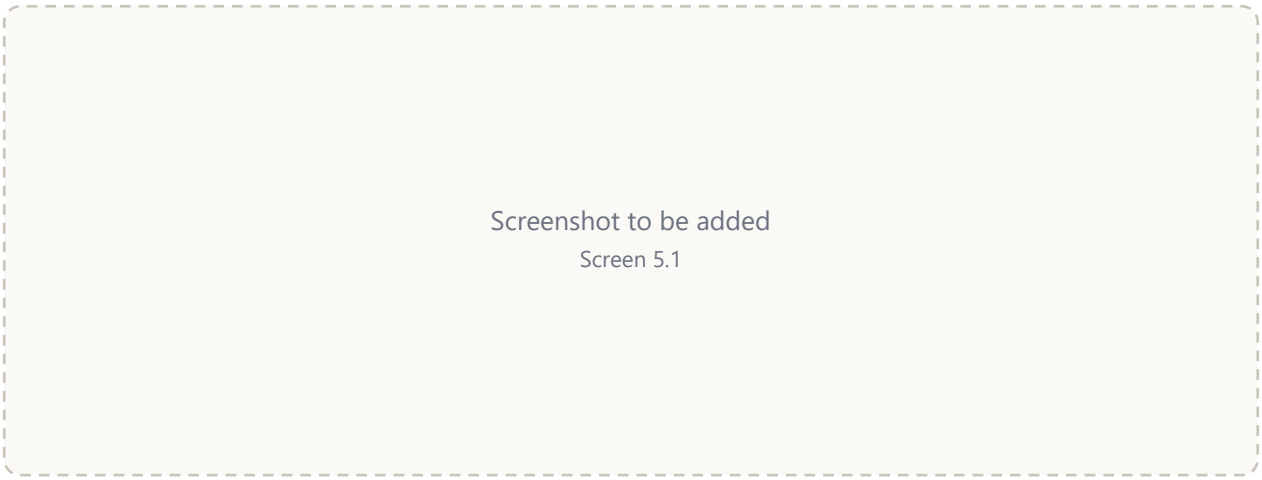
1. When an interview invitation arrives, open **Interview Management**.
2. Pick slots under **Free slots from your calendar**, then press **Draft reply with N slots** and approve the draft in **Pending Approvals**.
3. When the recruiter confirms, fill **Interview date**, **Interview time**, **Duration minutes** and **Meeting link**, then press **Confirm & add to calendar**.

If the calendar cannot be read, Cafy never assumes you are free: no slots are offered and the event is not created ("Calendar could not be checked; event not created."). Reconnect Google Calendar and try again. Outlook users get the same features through the Outlook connection.

SOP 5 — Career Analyst and Analytics

5.1 Ask the Career Analyst

1. Press the floating chat button (*Open Cafy Agent*) on any page.
2. Pick a quick action (*Analyse my opportunities*, *Explain a job match*, *Skill gap analysis*, *Application strategy*, *Show my progress*) or type a question.
3. Answers use only your résumé, current listings and your pipeline; job cards show **View job**. The chat needs the AI key on the server; without it the panel shows a service error.



Screen 5.1 — Career Analyst: quick actions (box), message box and send (arrow).

5.2 Analytics

1. Open **Analytics** (*Job Market & Skill Insights*): total stored jobs, active platforms, latest scrape results, most in-demand skills and platform distribution.
2. **Export Intel** downloads a CSV; **Sync All Crawlers** starts a fresh scrape of every job board.

SOP 6 — Guest user

Item	Guest	Signed-in user
How to start	Continue as Guest (when enabled)	SIGN UP / SIGN IN
Session length	12 hours, not renewable	Renews automatically (30 days)
Data kept	Private to that guest session	Kept in your account
Features	Same screens as a signed-in user	All features
Shown on screen	"Guest mode — use test data only"	Log Out in the sidebar

Use guest mode only with test data. Guest mode is on in local development and off in production unless the administrator sets **ENABLE_GUEST_MODE=true**.

SOP 7 — Administrator / Operator

Cafy Agent has no in-app admin role: every signed-in user sees the same screens. Administration is done on the server: installing, configuring, scheduling scrapes, deploying and backing up. Always use **uv** for Python packages, never pip.

7.1 Install and run locally

```
uv sync                                # Python 3.13 + locked packages
npm --prefix frontend install
uv run alembic -c backend/alembic.ini upgrade head   # Postgres only
uv run backend/run.py                  # API  http://127.0.0.1:8000  (docs: /docs)
npm --prefix frontend run dev          # UI   http://localhost:5173
uv add <package>                       # add a dependency (never pip install)
```

Windows: **backend\start.bat**; macOS/Linux: **backend/start.sh** (installs, frees port 8000, starts both). If uv reports *invalid peer certificate*, add **--system-certs** or set **UV_SYSTEM_CERTS=1**.

7.2 Configuration (environment variables)

Variable	Purpose	Production
DATABASE_URL	PostgreSQL URL; without it a local SQLite file is used	Required
JWT_SECRET_KEY	Session signing key (32+ characters)	Required
EMAIL_AGENT_SECRET	Signs email OAuth state; derives token encryption key	Required
TOKEN_ENCRYPTION_KEY	Encrypts stored mailbox tokens	Recommended
APP_ENV	production enables production checks	production
CRON_SECRET	Protects /api/cron and /api/email-agent/cron/poll	Recommended
PUBLIC_BASE_URL / FRONTEND_URL	OAuth redirect base / return page	Recommended
CORS_ORIGINS	Allowed browser origins (default *)	Set explicitly
GROQ_API_KEY, GROQ_MODEL	AI rewriting, Career Analyst, role summaries	Optional
NANGO_SECRET_KEY	Gmail and Google Calendar connections via Nango (integrations google-mail, google-calendar)	Optional
GOOGLE_CLIENT_ID/SECRET, MICROSOFT_CLIENT_ID/SECRET	Direct Gmail / Outlook OAuth	Optional
ENABLE_GUEST_MODE	Allow Continue as Guest	Off by default
SCRAPE_HOUR / SCRAPE_MINUTE / SCRAPE_TIMEZONE	Daily scrape time (default 10:00 Europe/London)	Optional
EMAIL_AGENT_POLL_MINUTES	Inbox check interval (default 5)	Optional
BRAVE_SEARCH_API_KEY	Web search for hiring-process research	Optional
CAFY_MATCH_CONFIG	JSON file overriding match-scoring settings	Optional

In production the API refuses to start without DATABASE_URL, a strong JWT_SECRET_KEY and EMAIL_AGENT_SECRET. Never commit .env; .env.example lists the main names.

7.3 Job collection (scraper)

1. Job boards are listed in **backend/scrapers/sources.yaml** (Greenhouse, Lever, Ashby, Workday, Workable, Oracle Cloud, BreezyHR, Pinpoint, SmartRecruiters). Only UK jobs are kept.
2. The scheduler runs once a day at SCRAPE_HOUR and catches up missed runs. On Vercel, the cron calls **/api/cron** (10:00 UTC) with **Authorization: Bearer CRON_SECRET**.
3. Run a scrape now: **Analytics → Sync All Crawlers**, or the API **POST /api/scrape?platform=all**.
4. Find more UK employers from the sponsor register: **uv run backend/scripts/discover_employers.py**.

7.4 Deploy

- **Docker**: **docker compose up -d** starts Postgres 18 and the API (migrations run on start; INSTALL_LATEX=true adds server-side PDF compilation). See docs/DOCKER.md.
- **Vercel**: push to main; set DATABASE_URL, JWT_SECRET_KEY, EMAIL_AGENT_SECRET, CRON_SECRET and GROQ_API_KEY. Background threads do not run there; the two daily crons replace them.
- After pulling changes run **npm --prefix frontend run build** so the served UI is current, then restart the API.

7.5 Backup and recovery

1. Back up: **docker/backup-local-db.ps1** (or .sh) writes a pg_dump file plus a .sha256 checksum to backups/.
2. Restore into Docker: **docker/restore-into-docker.ps1**; verify with **docker/compare_databases.py**.
3. Never run **docker compose down -v**; it deletes the database volume.
4. Smoke test a running server: **uv run docker/smoke_test.py --base URL**.

7.6 Monitoring and security

- Logs go to stdout (Docker: **docker compose logs -f api**). **GET /api/health** shows job count and scheduler state.
- Passwords use Argon2id; sessions are short JWTs plus rotating HttpOnly refresh cookies; failed logins are throttled.
- Mailbox tokens are encrypted at rest; Nango tokens are never stored by Cafy.
- Set CORS_ORIGINS and CRON_SECRET in production; rotate any secret that was ever committed.

How matching works

Find Jobs lists a job only when **all** of these hold. The same rule set is used on upload, refresh and by the Career Analyst.

Rule	Requirement
Experience level	Fresher (under 1 full-time year) or experienced, from dated full-time roles. A fresher never sees roles for experienced hires (senior/lead titles or 2+ years asked), and stated years must be met.
Score	85% or more.
Skills	At least 4 of the job's skills matched on the résumé.
Qualifications	Every essential qualification met: degree level, or named ones such as ACCA, CIPD, PRINCE2. "Or equivalent experience" is accepted for experienced candidates.
Not scored	Right to work, security clearance, driving licence, DBS: shown as checks before applying. Visa sponsorship is assessed separately.

How the score is calculated

1. Cafy reads the job description and groups its lines under headings: *Essential, Requirements, What we're looking for, About you…* are **essential**; *Desirable, Nice to have, Preferred, Bonus…* are **desirable**. Words on a line ("must", "ideally", "a plus") override the heading.
2. Each skill the advert names, in its own wording, is one requirement. Soft skills (e.g. communication) are listed but not scored.
3. Each requirement is graded against the whole résumé: **direct** (used in a role, internship, project or certification) 1.0; **listed only** (Skills list, summary, education) 0.5; **related tool** (e.g. Azure for AWS) 0.4; **missing** 0.
4. Essential requirements count twice a desirable one. **Score = points earned ÷ points possible**.
5. If the skills or qualification rule fails, the score is capped at 84% and the reason is recorded.

Example requirement	Importance	Résumé evidence	Credit × weight
Python	Essential	"Built Spark jobs in Python" (work)	1.0 × 2 = 2.0
Docker	Desirable	Listed in Skills only	0.5 × 1 = 0.5
AWS	Essential	Has Azure (same family)	0.4 × 2 = 0.8
Kafka	Desirable	Not on résumé	0 × 1 = 0

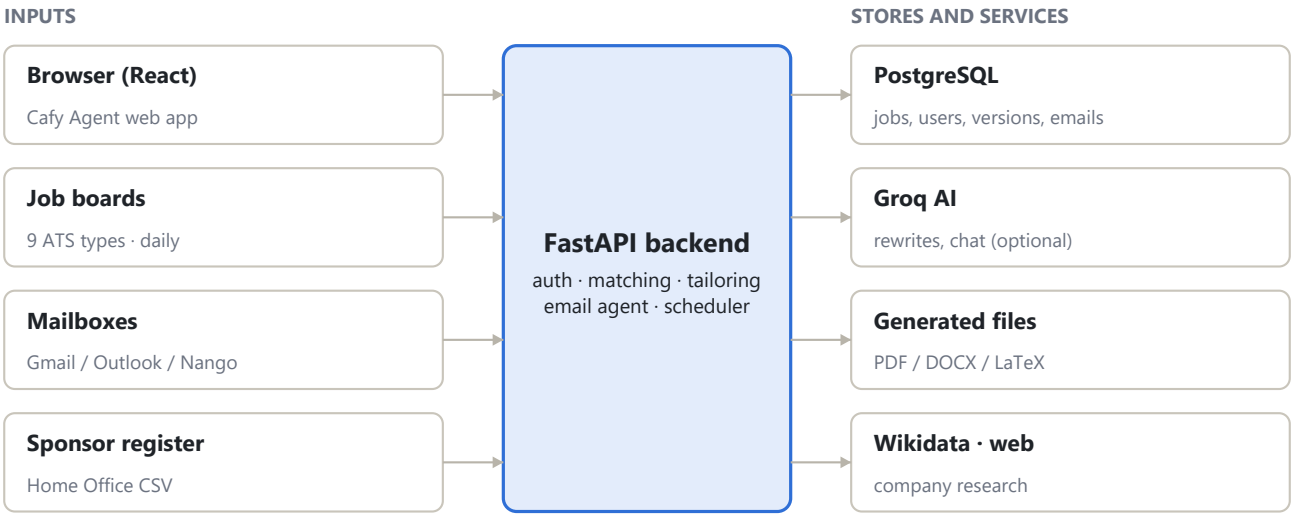
Weights are a transparent starting point, configurable with CAFY_MATCH_CONFIG and calibrated with backend/scripts/calibrate_matching.py against reviewed résumé–job pairs.

UK visa sponsorship

Employers are checked against the Home Office *Register of Worker and Temporary Worker licensed sponsors* (and Irish employment-permit lists for Ireland). **Visa Sponsorship** appears only when the employer is on the register **and** the job is confirmed in the UK. It is information, not a guarantee that a role offers sponsorship, and it never changes the match score.

System overview

How the parts connect



Component	What it does
Frontend (React, Vite)	All screens; renders the résumé PDF in the browser; served from frontend/dist
Backend (FastAPI, Python 3.13)	REST API under /api; per-user data isolation by account
Database	PostgreSQL in production (SQLite locally); schema via Alembic plus app tables
Scheduler	Daily job scrape; email monitor every 5 minutes (not on Vercel; crons instead)
AI provider	Groq (optional): résumé rewriting, Career Analyst, role summaries

Troubleshooting

Problem	Likely cause	What to do
Find Jobs shows no jobs after upload	No job meets the matching rules, or filters too narrow	Clear all filters; check your résumé lists your skills; wait for the next daily scrape
Skills missing from your profile	Skills heading not recognised or scanned PDF	Use a heading like “Skills”; upload a text-based PDF or DOCX
“Error processing résumé”	File over 10 MB, image-only PDF or unsupported type	Upload a smaller PDF/DOCX with selectable text
Generation failed	Job not found or server error	Retry; if it repeats, tell the administrator
AI Enhance / Career Analyst not working	AI key not configured	Administrator sets GROQ_API_KEY
Email Agent reads nothing	Mailbox not connected or token expired	Reconnect the mailbox; press Sync now; check the account error
“Nango is not configured”	NANGO_SECRET_KEY not set	Administrator sets NANGO_SECRET_KEY
No free slots / interview not added to calendar	Google Calendar not connected, or access expired	Connect or reconnect Google Calendar on the Email Agent page
“Too many failed attempts”	5 wrong passwords in 15 minutes	Wait the minutes shown, then sign in
Saved jobs disappeared	Pipeline is kept for the browser session only	Save again; download tailored résumés
Server will not start in production	Missing DATABASE_URL, JWT_SECRET_KEY or EMAIL_AGENT_SECRET	Set them; see 7.2

Known limitations

- **Job Alerts** page: shows “Job alerts are not available yet.”
- Saved jobs and statuses in **Applied To** are not stored on the server; a page reload clears them.
- □ **Add Job Manually** and the Staged Approvals filters are not functional yet.
- The detailed “Why this score?” breakdown is computed by the server but is not yet shown on a reachable screen.
- Certifications, publications and achievements are stored but not editable or printed in the current template.
- There is no in-app admin role or account deletion; administration is done on the server.
- The Account page says the scraper runs every 6 hours; it actually runs once a day (SCRAPE_HOUR).
- Outlook sync reads at most 50 messages per check.
- **Improve with AI** on a single section is instructed not to invent facts but is not automatically fact-checked (full tailoring is); review its result before keeping it.
- Skill detection tolerates word forms, so a résumé line like “account records” can count as *Accounting*; check skills added to a tailored résumé.
- Job alerts and alert emails are not wired up yet.

Appendix — main API endpoints

Method and path	Purpose
POST /api/auth/signup · /login · /refresh · /logout · /guest	Accounts and sessions
POST /api/resume/upload · GET /api/resume/current · DELETE /api/resume	Résumé and matches
GET /api/jobs · POST /api/recommend	Job list; matches (min_score default 85)
POST /api/resumes/generate · /save · /rewrite-section	Tailored résumé
GET /api/resumes/versions · DELETE /api/resumes/versions/{id}	Version history
GET /api/resumes/{id}/pdf · /docx · /latex	Downloads
GET /api/job-intelligence/{job_id}	Role, hiring process, match
POST /api/job-applications/start	Record an application
GET /api/email-agent/dashboard · POST …/accounts/{id}/sync · …/drafts/{id}/approve	Email Agent
POST /api/career-chat	Career Analyst
POST /api/export	Data export ZIP
GET /api/health · /api/stats · /api/matching/config	Status, analytics, scoring settings
GET /api/cron · /api/email-agent/cron/poll (Bearer CRON_SECRET)	Scheduled jobs

All /api routes except health, sign-in and cron need a signed-in session. Full interactive reference: **/docs** on the running API.